Simplified decision making in the belief space using belief sparsification

Vision-Aided Navigation -- 086761 -- Final Project

Guy Ben-Yosef	313580805
Noam Iluz	312163272

1. Introduction

In the realm of intelligent autonomous agents and robotics, the ability to make decisions under uncertainty is crucial, especially in applications like underwater exploration, space missions, and robotic automation. These decision-making tasks often involve complex, high-dimensional state spaces that require sophisticated reasoning about probabilistic states, referred to as "beliefs." A common framework to approach this is Belief Space Planning (BSP), often modeled as a Partially Observable Markov Decision Process (POMDP).

The paper titled "Simplified decision making in the belief space using belief sparsification," authored by Khen Elimelech and Prof. Vadim Indelman, focuses on improving the computational efficiency of BSP. The paper is of particular relevance to the field of aerospace engineering, especially for students and researchers at Technion's Faculty of Aerospace Engineering, as it aligns with the study of intelligent autonomous systems in aerospace applications.

1.1. Background and Importance

Autonomous agents like drones, robotic arms, and unmanned aerial vehicles (UAVs) need to function in uncertain environments, navigating and making decisions based on imprecise or incomplete data. In such contexts, decision-making involves optimizing actions to minimize the effects of uncertainty on the agent's ability to achieve its objectives.

Belief Space Planning is pivotal in scenarios where agents operate in partially observable environments, requiring them to plan their actions based on probabilistic beliefs about their state. The challenge lies in managing the high computational costs of reasoning with high-dimensional beliefs.

1.2. Relevance to the Course

The paper "Simplified decision making in the belief space using belief sparsification" aligns closely with the fundamental topics covered in the course, specifically vision-aided navigation (VAN) and simultaneous localization and mapping (SLAM). These subjects are essential for developing autonomous systems that can operate effectively in uncertain, dynamic, and unknown environments.

The course emphasizes Bayesian inference and state-of-the-art SLAM approaches, both of which rely heavily on belief space planning to handle uncertainties in the agent's environment. The paper's focus on simplifying decision-making by approximating the belief space directly contributes to improving the efficiency of belief space planning and active SLAM. This is particularly significant for aerospace students, as it directly impacts the development of robust, real-time autonomous navigation systems.

Additionally, the course deals with advanced topics like multi-robot cooperative localization and mapping and recent learning approaches. The paper's method, which reduces computational complexity in belief space planning, could potentially enhance multi-robot systems by enabling more efficient resource allocation and decision-making across agents. In the context of deep learning, the techniques discussed in the paper could also intersect with learning-based approaches to decision-making by providing structured approximations that can be incorporated into neural network models.

Overall, the paper's content is deeply relevant to the course's focus on cutting-edge methods in navigation and mapping, providing valuable insights and practical methods that can be directly applied in aerospace engineering.

1.3. Literature Review

Traditionally, solving BSP problems requires calculating an objective function for each candidate action, which becomes computationally expensive with high-dimensional belief spaces. The paper proposes a novel approach to simplify this problem by using belief sparsification, which approximates the belief state to retain key characteristics while reducing computational complexity.

Belief sparsification is not new in robotics, but most existing methods have focused on sparsifying beliefs for long-term operation and maintaining tractability rather than planning. This paper distinguishes itself by applying sparsification specifically for planning purposes, offering a framework that reduces computational complexity without compromising the accuracy of state inference.

1.4. Problem Importance

Efficient decision-making is essential for agents operating in dynamic and uncertain environments. The proposed belief sparsification technique offers a way to simplify decision-making by approximating the belief space. This simplification allows agents to make real-time decisions while ensuring reliable performance, which is critical for aerospace applications where quick and accurate decision-making is crucial.

The rest of the paper systematically introduces the belief sparsification technique, provides theoretical backing, and demonstrates its application in scenarios like active Simultaneous Localization and Mapping (SLAM), which is fundamental to autonomous navigation and robotic mapping in aerospace.

2. Preliminary material and problem formulation

Now, we wolud like to show the relevant notations and definitions that the article based on.

- State Variables (X): The agent's state, including its pose and other variables like landmarks.
- Belief (b): The agent's probability distribution over its state, denoted as $b(X)$.
- Information Matrix (Λ): The inverse of the covariance matrix of the belief.
- Collective Jacobian (U): The transformation applied to the state variables after executing control actions.
- Objective Function (V): Measures the agent's information gain or uncertainty reduction over a set of actions.

2.1. Mathematical Problem Definition

The problem tackled in the paper is formulated as:

2.1.1. Belief Propagation:

The belief b at time k, given previous controls $u_{1:k}$ and observation $z_{1:k}$ is:

$$b_k \;=\; P(X_k|u_{1:k}, z_{1:k}) \tag{1}$$

The belief is updated sequentially through control and observation models to maintain the posterior distribution over the state.

2.1.2. Decision-Making Objective:

The objective is to select a control sequence that minimizes uncertainty in the future belief, measured using the differential entropy of the posterior distribution:

$$H(b) \;=\; \frac{1}{2} ln\left[\frac{(2\pi e)^n}{|\Lambda|}\right] \;=\; -\frac{1}{2}(ln|\Lambda| - n \cdot ln(2\pi e)) \tag{2}$$

where Λ is the information matrix of the belief.The objective function for a candidate action u is:

$$V(b, u) \;=\; \frac{1}{2}\left(ln|\Lambda \;+ U^T U| - ln|\Lambda| - m \cdot ln(2\pi e)\right) \tag{3}$$

Here, U is the collective Jacobian of the action and m is the number of new variables introduced by the action.

2.2. Preliminary Mathematical Material

2.2.1. Information Matrix Updates:

Updates to the belief's information matrix Λ_k upon applying controls and observations are given by:

$$\Lambda_{k+1} = \Lambda_k + U^T U \tag{4}$$

where U is computed from the control and observation models.

2.2.2. Cholesky Factorization:

To compute the determinant efficiently, the paper relies on the Cholesky factorization:

$$\Lambda_k = R_k^T R_k \tag{5}$$

where R_k is an upper triangular matrix.

2.2.3. Entropy Calculation:

The entropy calculation involves the determinant of the information matrix.
The determinant is efficiently calculated using the diagonal elements of R_k from the Cholesky factorization.

3. Main contribution

The paper offers significant contributions to the field of robotics and autonomous systems, particularly in the realm of decision making under uncertainty. One of the major strengths of this work is the development of a scalable belief sparsification algorithm that facilitates efficient decision making by approximating high-dimensional belief states. This is particularly crucial in complex environments where computational resources and time are limiting factors. By introducing a method that maintains the consistency of solutions with the original, unsparsified problem, the authors address a common challenge in belief space planning and simultaneous localization and mapping tasks. Their approach not only reduces the computational load by simplifying the decision-making process but also ensures that the quality of the solutions is not compromised, thereby supporting the deployment of intelligent systems in real-time applications.

The experimental demonstration of the proposed belief sparsification method in an active-SLAM scenario marks a practical advancement. The results show a significant reduction in computation time without a loss in the quality of the solutions, which underscores the utility of the approach in real-world settings where autonomous agents must operate efficiently and reliably. This is particularly relevant in the field of computer vision aided navigation, where the ability to rapidly process and respond to dynamic environments can greatly enhance the performance of autonomous systems. Moreover, the theoretical framework established in this work provides a foundation for future research to build upon, offering a systematic way to quantify the effects of simplification on action selection. This could lead to further innovations in not just robotics, but also in areas where decision making under uncertainty is critical, such as autonomous vehicles and advanced surveillance systems.

The methodology introduced by the authors brings forth a nuanced approach to dealing with the inherent complexity of decision-making in high-dimensional belief spaces. Traditionally, planning in such environments required exhaustive calculations, encompassing every possible action and outcome, to ensure optimal decisions. This exhaustive approach is computationally prohibitive, particularly in real-time scenarios where decisions must be made swiftly to remain effective. By employing belief sparsification, the paper fundamentally shifts the paradigm by allowing a simplified yet effective examination of potential actions. This method systematically reduces the dimensionality of the belief space without sacrificing the accuracy of the decision-making process.

Mathematically, the core innovation lies in the formulation and implementation of the belief sparsification algorithm, which can be expressed through the algorithm's ability to maintain the entropy of the belief. The algorithm simplifies the decision problem while ensuring that the simplified belief b_s is a sparse approximation of the original belief b. Formally, this is represented by the equation

$$\Delta(P, P_s) = 0 \tag{6}$$

where P and P_s are the original and simplified decision problems respectively, indicating no loss in the quality of decision-making. This property is crucial as it underpins the theoretical guarantee that the action selected using the sparsified belief is as optimal as the one selected using the original belief. Additionally, the paper provides a mathematical framework to quantify the potential loss or offset in solution quality due to simplification, using a metric called the "simplification offset," $\delta(P, P_s, a)$, which measures the difference in the objective function values between the original and simplified problems across actions. This metric is vital for understanding the trade-offs involved in the simplification process and ensuring that they are within acceptable bounds for practical applications.

4. Discussion and critical thinking

One of the potential areas for further scrutiny within this paper is the assumption regarding the sparse approximation of the belief space and its practical applicability in highly dynamic environments. The authors have proposed a scalable belief sparsification algorithm which, while innovative, may face challenges in scenarios where the environment's dynamics rapidly change. In such cases, the sparsity induced by the algorithm might overlook some crucial but less apparent dependencies between state variables. This can lead to suboptimal decision-making or delayed responses to environmental changes. It would be beneficial for the authors to consider these dynamic scenarios more deeply, possibly integrating adaptive mechanisms that could re-evaluate and adjust the sparsity pattern as new data becomes available.

Another area that could be enhanced is the evaluation of the simplification's impact on decision-making quality in practical scenarios. The paper provides a theoretical framework and initial simulation results, which are promising. However, the transition from theoretical validation to real-world application often introduces non-idealities not accounted for in simulations. For instance, the noise characteristics in real sensors can deviate significantly from those modeled in simulations and can affect the reliability of the sparsification approach. Extending the validation to include real-world experiments or more complex simulation setups that better mimic actual operational conditions would not only strengthen the confidence in the results but also highlight potential limitations or necessary adjustments in the algorithm to accommodate various types of sensor inaccuracies and environmental unpredictability. This practical evaluation could provide a more comprehensive understanding of the utility and robustness of the proposed method.

The authors make several assumptions that could be considered optimistic, particularly in relation to the operational stability of the environment and the constancy of system dynamics. One assumption that stands out is the stability and predictability of the environmental conditions, which greatly influences the effectiveness of the belief sparsification strategy. In real-world scenarios, environments are often subject to sudden changes that are not always predictable, such as dynamic obstacles in robotics navigation or unexpected changes in terrain. The algorithm’s reliance on a priori sparsity patterns, which are assumed to be effective throughout the decision-making process, might not hold under such conditions. This could lead to decisions that are based on outdated or irrelevant information, thereby compromising the system’s performance and safety.

Furthermore, the assumption regarding the noise characteristics of sensors and their impact on the state estimation is another area where the paper’s assumptions may not align with practical applications. The authors base their sparsification technique on the assumption that noise in the system follows a known and somewhat stable statistical pattern, which allows for consistent sparsification of the belief space. However, in practical scenarios, sensor noise can be highly variable and dependent on numerous factors including environmental conditions, sensor wear and tear, and interference from other electronic devices. These factors can introduce non-linearities and non-Gaussian noise characteristics that are not accounted for in the model presented. Addressing these kinds of realistic sensor behaviors in the framework would enhance the applicability of the sparsification method and provide a more robust tool for real-world decision-making scenarios.

A promising direction for future research stemming from this paper could be the development of adaptive sparsification algorithms that dynamically adjust to changes in the environment or system state. Current methods, including those discussed in the paper, generally rely on a static sparsification pattern established based on the initial understanding of the environment. An adaptive approach could continuously or periodically re-evaluate the sparsity pattern based on real-time data, accommodating unexpected environmental changes, non-stationary noise characteristics, and other dynamics that may affect decision quality. This could involve the integration of machine learning techniques, such as reinforcement learning or online learning algorithms, that learn the optimal sparsity level and structure as they interact with the environment. Such advancements would make the sparsification process more robust and applicable across a wider range of scenarios, particularly in fields like autonomous driving and mobile robotics, where environmental unpredictability is a major challenge.

Another potential area for further investigation could be the exploration of hybrid models that combine sparse approximations with full belief updates at critical moments. This hybrid approach could allow the system to maintain computational efficiency through sparsification while periodically recalibrating with a full belief update to correct any drift or error accumulation introduced by the sparsification. Implementing such a strategy could be extremely beneficial in high-stakes environments such as surgical robotics, where precision is crucial and the stakes are high, or in autonomous vehicle navigation through densely populated urban areas, where dynamic obstacles and unpredictable events frequently occur. In these scenarios, the cost of integrating occasional comprehensive updates is far outweighed by the potential risk of inaccurate or outdated information leading to failures or accidents. Research could focus on optimizing the timing and frequency of these full updates to maximize safety without compromising the system's efficiency.

5. Our contribution

As been said, the original paper assumes a static sparsification approach that does not adapt to changes in the environment or the system's state. The sparsification is determined once per planning cycle and does not consider subsequent environmental interactions that might influence the importance of certain states or variables.

To make the sparsification process more dynamic and responsive, we propose an adaptive belief sparsification mechanism. This mechanism would adjust the sparsity pattern based on feedback from the environment or changes in the system's state.

Let ξ be the initial belief state, and ξ_s be the sparsified state. In the adaptive framework, the sparsity pattern S is updated based on a feedback function f that measures the deviation between the predicted outcomes and the actual outcomes observed after applying the chosen actions. The Sparsity Pattern Update:

$$S_{t+1} = g(S_t, f(\xi_{s,t}, \xi_{obs,t})) \tag{7}$$

where

S_t is the sparsity pattern at time t.

$\xi_{s,t}$ is the sparsified belief state at time t used for decision making.

$\xi_{obs,t}$ is the observed state after executing actions, reflecting the real environment feedback.

g is an update function that determines how the sparsity pattern should be adjusted based on the feedback.

This approach allows the system to dynamically adjust its belief sparsity, potentially leading to more efficient and accurate decision-making under varying conditions.
The paper mainly focuses on Gaussian beliefs, which might not adequately represent systems where beliefs are inherently multi-modal, such as in scenarios involving ambiguous or non-linear dynamics. We propose to extend the sparsification technique to support multi-modal distributions. This could involve using a mixture of Gaussians or particle filters to represent the belief state, providing a richer representation that captures multiple hypotheses about the system's state.
Consider a belief represented as a mixture of Gaussians:

$$b_k = \sum_{i=1}^{N} \omega_i \mathcal{N}(\mu_i, \Sigma_i) \tag{8}$$

Where ω_i are the weights, and $\mathcal{N}(\mu_i, \Sigma_i)$ are the Gaussian components of the mixture. The Sparsified Multi-modal Belief:

$$b_{s,k} = \sum_{i=1}^{N} \omega_i \mathcal{N}(\mu_i, \Sigma_{s,i}) \tag{9}$$

Where $\Sigma_{s,i}$ are the sparsified covariance matrices obtained by applying a sparsification matrix M that optimizes the Kullback–Leibler divergence[1] between the original and the sparsified distributions. The sparsification matrix M is selected to minimize the information loss across all components:

$$M = \mathrm{argmin}_M \sum_{i=1}^{N} \omega_i D_{KL}\left(\mathcal{N}(\mu_i, \Sigma_i) \,||\, \mathcal{N}\left(\mu_i, M\Sigma_i M^\top\right)\right) \tag{10}$$

where D_{KL} denotes the Kullback–Leibler divergence.
The Kullback–Leibler (KL) divergence is a statistical measure used to quantify the difference between two probability distributions. In the context of belief sparsification, the KL divergence serves as a critical tool for evaluating how well the sparsified belief state approximates the original belief state. By minimizing the KL divergence, we aim to retain the most crucial information in the belief state while reducing its complexity. This is especially important in dynamic environments, where the ability to make accurate decisions quickly can significantly impact system performance. In our adaptive sparsification framework, we utilize the KL

[1] Kullback, S., & Leibler, R. A. (1951). On information and sufficiency. The annals of mathematical statistics, 22(1), 79-86.

divergence to determine the optimal sparsification matrix M. This matrix is designed to minimize the information loss when the original Gaussian components of the belief state are approximated by their sparser versions. Consequently, maintaining a low KL divergence ensures that the essential characteristics of the belief state are preserved, thereby enabling more reliable decision-making under uncertainty.

These extensions aim to make the decision-making process more robust and adaptable to complex, dynamic environments. Future work could involve developing algorithms for optimal dynamic sparsity pattern adjustment and evaluating the performance of the adaptive and multi-modal sparsification frameworks in real-world scenarios, such as robotics navigation, autonomous driving, and other AI systems operating in uncertain environments.

6. Conclusions

In this report, we explored the concept of belief sparsification as presented in the paper "Simplified decision making in the belief space using belief sparsification" by Khen Elimelech and Prof. Vadim Indelman. The paper offers a significant contribution to the field of robotics and autonomous systems, specifically in the domain of decision-making under uncertainty. By introducing a scalable belief sparsification algorithm, the authors provide a method to reduce computational complexity without compromising the accuracy of state inference, which is particularly crucial in high-dimensional belief spaces.

6.1. Key Takeaways

The belief sparsification method is highly relevant for applications such as underwater exploration, space missions, and robotic automation, where agents must navigate and make decisions under uncertainty. This method aligns well with the core topics of vision-aided navigation (VAN) and simultaneous localization and mapping (SLAM) covered in our course, emphasizing the practical importance of reducing computational load while maintaining decision accuracy. The development of a belief sparsification algorithm that simplifies the belief state representation makes decision-making more computationally efficient. The authors demonstrated the algorithm’s effectiveness in active SLAM scenarios, showing significant reductions in computation time without loss of solution quality.

While the method shows promise in reducing computational complexity, it does not specifically address scenarios with highly dynamic environments. This could potentially lead to suboptimal decision-making or delayed responses to rapid environmental changes. To fully understand the method's applicability and limitations in practical scenarios, further validation through real-world experiments or more complex simulations is necessary.

6.2. Our Contribution

We proposed an enhancement to the belief sparsification method to address its limitations in dynamic environments. Our adaptive sparsification mechanism dynamically adjusts the sparsity pattern based on feedback from environmental changes, improving the system's responsiveness and decision accuracy. This involves developing a dynamic detection algorithm to identify significant changes in the environment and formulating an adaptive sparsification algorithm that modifies the belief state representation based on detected changes. Additionally, we introduced a sparsity control parameter that adjusts the degree of sparsity as a function of the variance in sensor data changes. We also extended the method to support multi-modal distributions, providing a richer representation that captures multiple hypotheses about the system's state by representing beliefs as a mixture of Gaussians and applying sparsification to each Gaussian component while minimizing the Kullback–Leibler divergence between the original and sparsified distributions.